

PROJECT REPORT ON

TrustGuard: Data Quality Pipeline for Retail Data

SUBMITTED BY

Raghav

Bachelors of Technology (B.Tech)

Computer Science and Engineering

Maharishi Markandeshwar (Deemed To Be University)

Intern ID - CT_CSI_DE_104

University Roll No.: 11232785

SUBMITTED TO

Celebal Technologies

(Data Engineering)

1. Introduction

The **TrustGuard: Data Quality Pipeline for Retail Data** project is an end-to-end Data Engineering project developed using **Apache Spark, PySpark, Delta Lake, SQL, and Databricks**. The primary objective of this project is to process raw retail transaction datasets, perform comprehensive data quality checks, clean inconsistent records, and generate meaningful business insights that help retail organizations make accurate and data-driven decisions.

In today's digital retail environment, supermarkets, e-commerce platforms, and retail stores generate a massive amount of transaction data every day. This data includes customer information, product categories, purchased items, quantities, prices, payment methods, transaction dates, and store locations. Processing such large volumes of raw data manually is difficult, time-consuming, and prone to errors. Therefore, Data Engineering techniques are used to ingest, validate, clean, and transform raw data into a structured and analytics-ready format.

This project follows a modular Data Engineering pipeline consisting of five major stages: **Data Ingestion, Data Quality Checks, Data Cleaning, SQL-Based Analysis, and Basic Anomaly Detection**. The Data Ingestion module loads raw CSV files into Databricks, while the Data Quality Checks module identifies missing values, duplicate records, invalid formats, and other data quality issues. The Data Cleaning module standardizes and transforms the data into a consistent format suitable for further analysis. SQL-Based Analysis generates business reports such as monthly revenue, top customers, payment method distribution, and product category performance. Finally, the Basic Anomaly Detection module identifies unusual records, such as invalid prices or abnormal transaction quantities, and stores them for further investigation without interrupting the pipeline.

The project uses a retail transaction dataset containing information such as **Transaction ID, Customer ID, Category, Item, Price per Unit, Quantity, Total Spent, Payment Method, Location, Transaction Date, and Discount Applied**. The dataset is processed using Apache Spark to ensure scalability and efficiency while handling large amounts of data. Cleaned datasets are stored as **Delta Tables**, providing reliable storage, and improved query performance.

Several Data Engineering operations are performed throughout the project, including data ingestion, schema validation, null value handling, duplicate removal, data standardization, type casting, SQL analytics, and anomaly detection. Business insights are generated using SQL queries to analyze customer spending patterns, monthly revenue, payment method usage, and product category performance. These insights help retail organizations monitor sales trends, understand customer behavior, and improve business strategies.

Overall, this project demonstrates the practical implementation of a complete **ETL (Extract, Transform, Load)** pipeline using modern Data Engineering tools and technologies. It provides hands-on experience in building scalable data pipelines, improving data quality, generating analytical reports, and implementing anomaly detection techniques. The project also highlights the importance of high-quality data in supporting reliable business intelligence and effective decision-making in the retail industry.

2. Objective

The primary objective of this project is to design and implement a complete **Data Quality Pipeline** using Apache Spark, PySpark, SQL, Delta Lake, and Databricks. The project aims to demonstrate how raw retail transaction data can be ingested, validated, cleaned, transformed, and analyzed to generate meaningful business insights. Through this project, the complete lifecycle of data processing is implemented, starting from raw CSV files and ending with clean, analytics-ready datasets stored as Delta Tables.

Another important objective of this project is to gain practical exposure to **ETL (Extract, Transform, Load)** concepts that are widely used in modern Data Engineering. The project focuses on handling structured retail datasets efficiently, improving data quality, and preparing reliable data for business analysis. It also helps in understanding how organizations build automated data pipelines to convert raw operational data into valuable business information.

The project is designed not only to perform technical Data Engineering operations but also to solve real-world retail data quality challenges. By identifying missing values, duplicate records, inconsistent formats, invalid entries, and unusual transaction patterns, the project improves the reliability of business data. In addition, SQL-based analytical reports provide valuable insights into customer spending, monthly revenue, payment methods, and product category performance, helping retail organizations make informed business decisions.

The main objectives of this project are:

- To read and process retail transaction CSV datasets using Apache Spark and PySpark.
- To implement a structured Data Quality Pipeline using Databricks.
- To perform data ingestion and schema validation.
- To identify and handle missing values, duplicate records, and inconsistent data.
- To standardize and transform raw data into a clean and analytics-ready dataset.
- To perform SQL-based business analysis on cleaned retail data.
- To generate business insights such as monthly revenue, top customers, payment method distribution, and product category performance.
- To implement basic anomaly detection using statistical threshold rules.
- To store processed datasets as Delta Tables for reliable and efficient data management.
- To gain hands-on experience with ETL (Extract, Transform, Load) pipeline development using modern Data Engineering tools.
- To understand the importance of data quality in improving business intelligence and decision-making.
- To build a beginner-friendly Data Engineering project that demonstrates practical implementation of Apache Spark, Delta Lake, SQL, and Databricks.
- To prepare reliable, clean, and scalable datasets that can be used for reporting, analytics, and future business applications.

3. Tools and Technologies Used

The successful implementation of this project required the use of several modern Data Engineering tools and technologies, each serving a specific purpose throughout the data quality pipeline. These technologies were selected to efficiently perform data ingestion, validation, cleaning, transformation, SQL-based analysis, and anomaly detection while following industry-standard Data Engineering practices. Apache Spark and PySpark were used as the core processing framework, whereas Python provided the programming environment for implementing the ETL pipeline. Databricks served as the cloud-based platform for developing and executing the pipeline, while Delta Lake was used to store processed datasets with ACID transaction support. SQL was used to generate business insights, and CSV files served as the primary data source. Together, these technologies enabled the development of a scalable, reliable, and efficient data quality pipeline.

Tools and Technologies Used:

- Python 3.x
- Apache Spark
- PySpark
- Databricks
- Delta Lake
- SQL

4. Dataset Used

The project uses a **Retail Transactions Dataset** containing customer purchase information collected from retail business operations. The dataset represents day-to-day sales transactions and includes customer details, product information, payment methods, transaction dates, and store locations. Before analysis, the dataset was loaded into Databricks using Apache Spark and processed through different stages of the data quality pipeline, including data ingestion, validation, cleaning, SQL-based analysis, and anomaly detection.

The dataset provides sufficient information to perform data quality checks, generate business insights, and detect unusual transaction patterns. After cleaning and validation, the processed dataset is stored as Delta Tables for efficient querying and further analysis.

4.1 Retail Transactions Dataset

The Retail Transactions Dataset contains detailed information about customer purchases made at retail stores. It is the primary dataset used throughout the project for implementing the complete Data Quality Pipeline. The dataset includes the following attributes:

- Transaction ID
- Customer ID
- Category
- Item
- Price Per Unit
- Quantity

- Total Spent
- Payment Method
- Location
- Transaction Date
- Discount Applied

The dataset is used to perform data ingestion, schema validation, null value handling, duplicate removal, data standardization, SQL-based business analysis, and anomaly detection. It also helps generate valuable business insights such as monthly revenue, top customers, payment method distribution, and product category performance.

5. Project Workflow

The **TrustGuard: Data Quality Pipeline for Retail Data** project follows a structured Data Engineering workflow to process raw retail transaction data into clean, reliable, and analytics-ready datasets. Instead of directly analyzing the raw CSV file, the project processes the data through multiple stages such as data ingestion, validation, cleaning, SQL-based analysis, and anomaly detection. This structured workflow ensures that the data remains accurate, consistent, and suitable for business reporting.

The project is implemented as a modular data pipeline consisting of five major modules. Each module performs a specific task before passing the processed data to the next stage. This approach improves maintainability, simplifies debugging, and follows modern Data Engineering practices.

During the implementation of this project, several Data Engineering operations were performed, including loading CSV data into Databricks, validating the dataset structure, identifying missing values and duplicate records, cleaning inconsistent data, performing SQL-based business analysis, and detecting anomalous transactions. The processed data was stored as **Delta Tables**, providing reliable storage, improved query performance, and ACID transaction support.

The workflow was executed step-by-step to ensure that each stage produced validated and reliable data before moving to the next module. This approach closely resembles real-world data engineering pipelines used by organizations to process operational retail data. The complete workflow implemented in this project is described below.

Step 1 – Data Ingestion

The raw retail transaction CSV file is loaded into Databricks using Apache Spark. The pipeline validates the dataset schema and checks whether all expected columns are present before storing the raw data in the **Raw Layer**.

Step 2 – Data Quality Checks

The raw dataset is analyzed to identify common data quality issues such as missing values, duplicate records, invalid formats, and schema inconsistencies. A Data Quality Report is generated to summarize the condition of the incoming dataset before cleaning begins.

Step 3 – Data Cleaning

The identified data quality issues are resolved by handling missing values, removing duplicate records, standardizing text and date formats, performing type casting, and validating important fields. Invalid records are separated whenever necessary, while valid records are transformed into a clean and analytics-ready dataset. The cleaned dataset is stored as a **Delta Table** in the **Clean Layer**.

Step 4 – SQL-Based Analysis

SQL queries are executed on the cleaned dataset to generate meaningful business insights. Reports such as monthly revenue analysis, top customers by total spending, payment method distribution, and product category performance are generated and stored as Delta Tables for reporting and business analysis.

Step 5 – Basic Anomaly Detection

The cleaned dataset is analyzed using simple statistical threshold rules to detect unusual transaction records. Transactions with abnormal quantities or invalid unit prices are flagged and stored in an **Anomaly Log** along with the reason for detection. This module runs after data cleaning and does not interrupt the overall pipeline

6. Results

The implementation of the **TrustGuard: Data Quality Pipeline for Retail Data** project was successfully completed using Apache Spark, PySpark, SQL, Delta Lake, and Databricks. The project successfully ingested the raw retail transaction dataset, performed comprehensive data quality checks, cleaned inconsistent records, generated business insights, and detected anomalous transactions using a structured data engineering pipeline.

Initially, the retail transaction dataset was successfully loaded into Databricks and stored in the **Raw Layer** without any modifications. The dataset was then validated through multiple data quality checks to identify missing values, duplicate records, inconsistent formats, invalid entries, and schema mismatches. These quality checks provided a clear understanding of the condition of the incoming data before any transformations were applied.

After completing the validation process, various data cleaning operations were performed, including handling missing values, removing duplicate transaction records, standardizing text values, formatting transaction dates, performing data type conversions, and validating important fields. The cleaned dataset was then stored as a Delta Table in the Clean Layer, ensuring reliable storage and improved query performance.

In the next stage of the project, SQL-based analytical reports were generated from the cleaned dataset. These reports provide meaningful business insights that can help retail organizations monitor sales performance, understand customer purchasing behavior, evaluate payment method usage, and analyze product category performance. The analytical reports were also stored as Delta Tables for future reporting and business analysis.

The final module implemented **Basic Anomaly Detection** using simple statistical threshold rules. Transactions with unusually high quantities and invalid unit prices were identified and stored in an **Anomaly Log** along with the corresponding reason for detection. This process enables organizations to identify suspicious or abnormal transactions without interrupting the normal data processing pipeline.

The project successfully generated several valuable business insights:

- Monthly revenue analysis based on transaction data.
- Top 10 customers based on total spending.
- Payment method distribution and usage analysis.
- Product category performance based on revenue and transaction count.
- Data Quality Report showing missing values and duplicate records.
- Basic anomaly detection using statistical threshold rules.
- Clean and analytics-ready retail dataset stored as Delta Tables.

The successful implementation of this project demonstrates how modern Data Engineering techniques can improve data quality, enhance reporting accuracy, and support better business decision-making. The pipeline provides a scalable and reliable solution for processing retail transaction data while ensuring that only high-quality data is used for analytical purposes.

7. Conclusion

This project successfully demonstrates the implementation of a complete **Data Quality Pipeline** using **Apache Spark, PySpark, SQL, Delta Lake, and Databricks**. The project provided practical experience in processing raw retail transaction datasets and transforming them into clean, reliable, and business-ready data through a structured **ETL (Extract, Transform, Load)** workflow. By implementing different stages of data processing, the project helped in understanding how raw retail data can be converted into meaningful business information for analysis and decision-making.

Throughout the project, various Data Engineering concepts were implemented, including data ingestion, schema validation, data quality checks, data cleaning, data transformation, SQL-based analysis, Delta Table creation, and anomaly detection. The project also demonstrated the importance of maintaining high data quality by identifying missing values, duplicate records, invalid entries, and inconsistent data before performing analytical operations. These processes ensured that the final dataset was accurate, consistent, and suitable for business reporting.

The modular pipeline architecture played a significant role in organizing the overall workflow. Each module—Data Ingestion, Data Quality Checks, Data Cleaning, SQL-Based Analysis, and Basic Anomaly Detection—performed a specific responsibility within the pipeline. This modular approach improved the maintainability, scalability, and readability of the project while following industry-inspired Data Engineering practices.

The project successfully generated several valuable business insights, including monthly revenue analysis, top customer analysis, payment method distribution, and product category performance. These analytical reports can help retail organizations improve reporting accuracy, monitor sales trends, understand customer purchasing behavior, and support data-driven business decisions. Overall, this project provides a strong foundation in modern Data Engineering concepts and demonstrates the practical implementation of a scalable and reliable retail data quality pipeline.

8. Future Scope

Although the TrustGuard: Data Quality Pipeline for Retail Data project successfully implements a complete Data Engineering pipeline, there is scope for further enhancement. The current project processes a static retail transaction dataset, but it can be extended to handle real-time data from retail stores and online platforms using technologies such as Apache Kafka or Spark Structured Streaming.

The project can also be improved by integrating machine learning-based anomaly detection for more accurate identification of unusual transactions. Interactive dashboards can be created using **Power BI** to visualize business insights, and the pipeline can be deployed on cloud platforms such as **Microsoft Azure** or **AWS** for better scalability and performance.

The project can be further enhanced by:

- Integrating real-time data streaming using Apache Kafka or Spark Structured Streaming.
- Developing interactive dashboards using Power BI or Tableau.
- Deploying the pipeline on cloud platforms such as Microsoft Azure or AWS.
- Implementing advanced anomaly detection and automated data quality validation.